

INDEPENDENT PRODUCTION AUDIT

UrbanGrid SEO, GEO & LLM Visibility Audit

Page speed, metadata, schema, security, AI crawler access, content citability, and current search-grounding visibility.

Website: <https://urbangrid.ae>

Date: 9 September 2026

Evidence: Live production checks, source review, three security scanners, three specialist reviews, and current discovery queries.

Executive scorecard

Engine	Score	Assessment
SEO Engine	62/100	Sound crawl basics and metadata, weakened by SPA rendering, duplicate H1s, unsupported schema claims, broken social imagery, weak caching, and security risk.
GEO Engine	55/100	AI crawlers are allowed and service/FAQ schema exists, but llms.txt is invalid, core copy is JS-dependent, citations are sparse, and entity evidence is inconsistent.
LLM Visibility	45/100	Discoverable for branded searches, absent from the sampled top ten for two high-intent non-branded query groups.
Security	Critical	A public API exposes inspector records including password-hash fields, and privileged credentials exist in source.

Strong foundation: HTTPS, HTTP/2, crawler access, robots.txt, sitemap, canonical tags, server-injected metadata, schema, and healthy measured TTFB.

Immediate priority: Security remediation must precede incremental metadata work.

Critical findings

P0 — Public inspector data and password-hash exposure

An unauthenticated production request to `/api/inspectors` returned HTTP 200 with inspector records containing username, passwordHash, full name, email, and phone fields.

Action: Require admin authentication, return a restricted DTO, rotate affected credentials, and review access logs.

P0 — Hardcoded privileged credentials

Privileged admin and inspector credentials are embedded in server source. Values are intentionally omitted.

Action: Remove and rotate them, disable automatic default accounts, enforce managed secrets, and add MFA or managed identity.

P1 — 115 dependency vulnerabilities

2 critical, 46 high, 53 moderate, 14 low. Critical findings affect fast-xml-parser and have non-major fixes available. High findings include axios, drizzle-orm, nodemailer, rollup, ws, and postcss.

P1 — Invalid llms.txt

`/llms.txt` returns HTTP 200 with `text/html` and the homepage shell, not a plain-text LLM resource.

SEO Engine

What works

- HTTPS, HTTP/2, HSTS, CSP, nosniff, clickjacking protection, referrer policy, and permissions policy.
- Valid robots.txt with sitemap reference; sitemap covers core, service, location, and blog pages.
- Route-specific titles, descriptions, canonicals, and JSON-LD are present in tested initial HTML.
- Measured fresh TTFB was approximately 178 ms.

P1 — JavaScript-dependent body content

The server injects metadata, schema, and a hidden H1, but most useful copy and links arrive after React hydration. Pre-render the primary answer, evidence, FAQs, and internal links.

P1 — Duplicate H1 signals

Tested HTML contains a route-specific hidden H1 plus the generic hero H1. Use one visible, route-specific H1.

P1 — Inconsistent trust claims

Claims such as 40,000+ inspections, 600,000+ defects, 2,000+ clients, 10+ years, and review scores need consistent definitions and evidence.

P2 — Unsupported ratings and misleading location schema

Aggregate ratings lack visible supporting reviews. Emirate pages reuse a Dubai address. Remove unsupported ratings and distinguish genuine offices from service areas.

P2 — Broken social image and weak asset caching

The declared OG image resolves as HTML. Fingerprinted assets use `private, max-age=0`. Main JS is 403,503 bytes raw and main CSS is 204,329 bytes raw. Publish the image, add immutable caching, and reduce hydration cost.

GEO Engine

AI crawler access: strong

GPTBot, ChatGPT-User, ClaudeBot, Google-Extended, PerplexityBot, Bingbot, and Googlebot all received HTTP 200.

Citability: weak-to-moderate

- Marketing claims often lack named sources, dates, or methodology.
- Case studies need property details, tools, defect categories, totals, and outcomes.
- Expert bylines, reviewer credentials, and authoritative sameAs profiles are limited.
- Generic hero content reduces page-specific answer density.

Priority improvements

- 1. Publish a valid llms.txt and a canonical content inventory.
- 2. Lead priority pages with a 40–60 word direct answer.
- 3. Use question-led headings, comparison tables, checklists, and concise FAQs.
- 4. Cite named UAE regulators, standards bodies, dated sources, and expert reviewers.
- 5. Publish evidence-rich inspection case studies and original datasets.

LLM Visibility Engine

Current query group	Observed result
Best property snagging / building inspection companies in UAE and Dubai	Absent from sampled top 10
Building condition assessment / technical due diligence providers in UAE	Absent from sampled top 10
What is UrbanGrid UAE / branded inspection query	Present, approximately position 3

UrbanGrid has branded entity recognition but weak non-branded category visibility. One indexed careers snippet still used “Lena,” showing stale external indexing.

Method limitation: This run used current web-grounding evidence, not authenticated consumer/API queries to ChatGPT, Claude, and Gemini. A direct share-of-voice benchmark should run the same 20–30 prompts across all three platforms from the target geography and repeat monthly.

Security detail

Strengths

- Strong baseline HTTP security headers.
- Automated SAST returned no findings.
- HoundDog returned five privacy/dataflow findings: one medium and four low.

Confirmed risks

- Public inspector endpoint exposes PII and password-hash fields.
- Hardcoded privileged credentials and a fallback session secret.
- Session cookies explicitly configured with `secure: false`.
- API response bodies and contact PII can be written to logs.
- Unauthenticated conversation/model routes exist in source; live registration was not confirmed.

Prioritized action plan

First 24 hours

1. Protect `/api/inspectors` and remove sensitive fields.
2. Rotate and remove hardcoded credentials.
3. Enforce the production session secret and secure cookie settings.
4. Review endpoint access logs.

First 7 days

1. Patch critical/high runtime dependencies and rescan.
2. Publish `llms.txt`.
3. Fix OG imagery.
4. Remove unsupported ratings.
5. Reconcile company claims.
6. Stop logging PII and API bodies.

First 30 days

1. Pre-render priority page content.
2. Correct H1 and location schema.
3. Add citable evidence, expert bylines, and dates.
4. Add immutable asset caching.
5. Launch a 20-30 prompt LLM benchmark.

60-90 days

1. Publish original UAE inspection data and case studies.
2. Earn relevant third-party citations.

3. Strengthen verified entity/NAP profiles.
4. Track non-branded ranking, AI mentions, citations, and leads monthly.

Confidence and scope

High confidence: Live headers, crawler responses, robots/sitemap, invalid llms.txt, exposed inspector endpoint, hardcoded credentials, schema, asset caching, and bundle sizes.

Medium confidence: Composite SEO/GEO scores.

Limited: Direct proprietary-model visibility and field Core Web Vitals. PageSpeed API quota was unavailable, so no Lighthouse score was invented.

This is a point-in-time technical and visibility audit, not a guarantee of rankings, citations, or security. Findings should be revalidated after remediation.